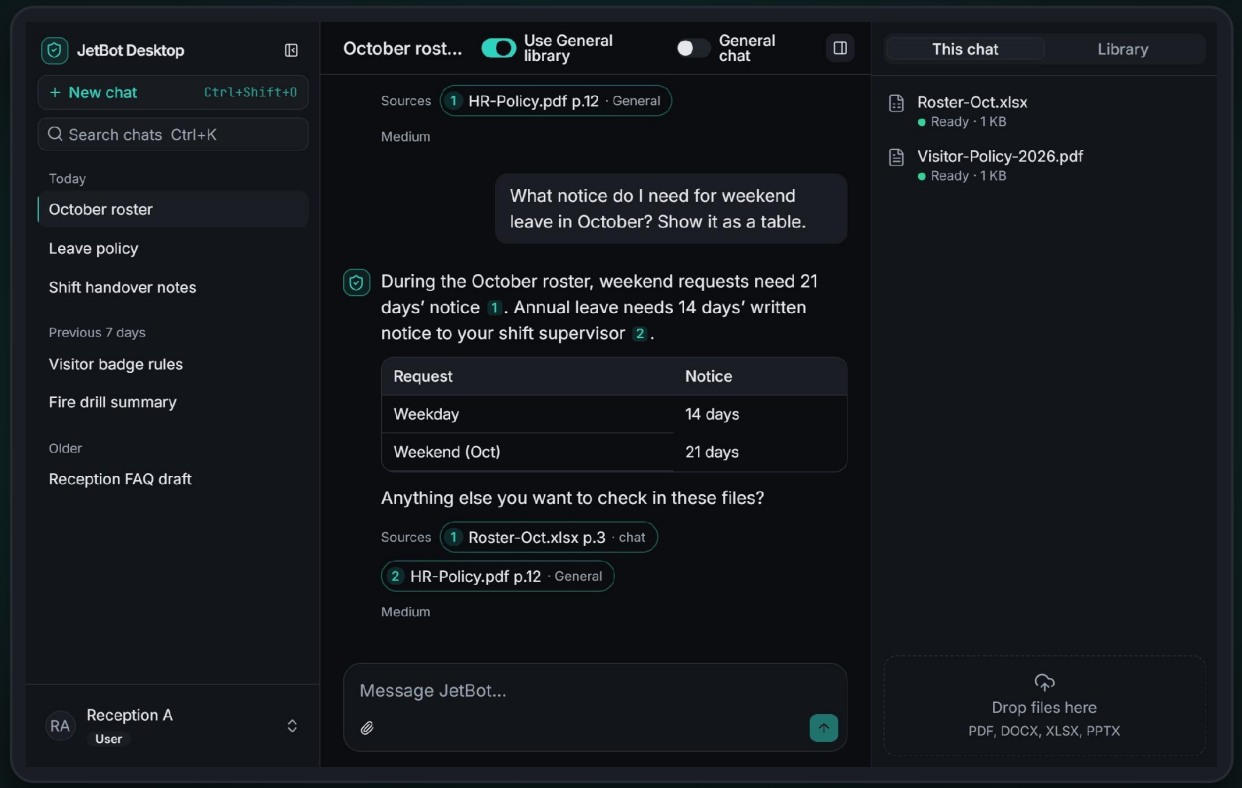
JetBot Desktop

Private AI for your documents

An offline assistant that answers questions from your own files, with citations, on a PC you already own. Nothing leaves the building.



Product overview · Version 0.1.0 · October 2026

THE PROBLEM

Teams want AI. Their documents cannot go to the cloud.

Regulated and sensitive work

Defence, government, healthcare, legal and finance teams are barred from pasting documents into public AI services.

Restricted or air-gapped sites

Many sites have no internet at all, or only a filtered one. Cloud AI simply does not reach them.

Per-seat cloud costs add up

Enterprise AI subscriptions are priced per user per month, and shared shift PCs do not fit that model.

Knowledge is locked in files

Policies, procedures and manuals sit in PDFs and Word files that staff search by hand during busy shifts.

THE SOLUTION

Cited answers from your files, on one offline PC

100% offline

Model, search index, database and files all live on the machine.

Grounded and cited

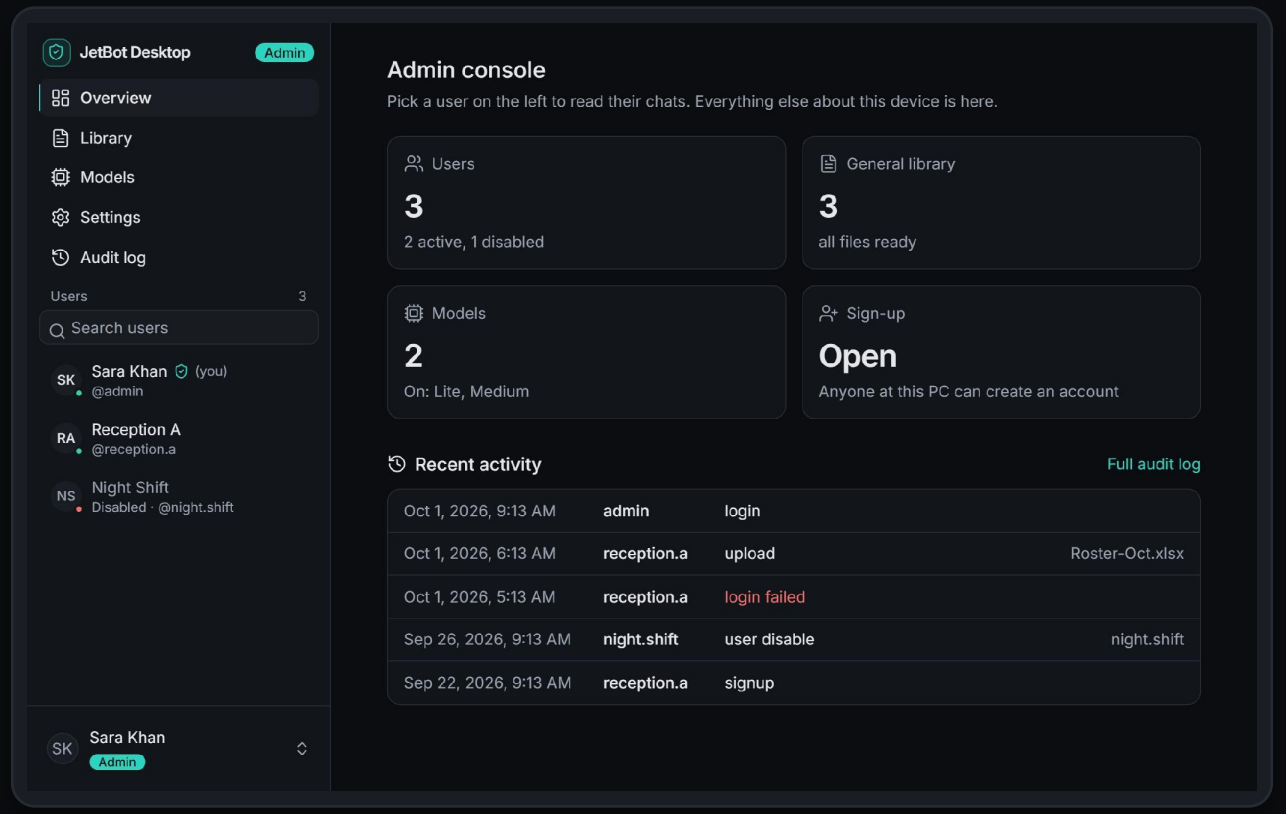
Every answer names the file and page it came from.

Built for shared PCs

Each person signs in; their chats and files stay private.

Admin in control

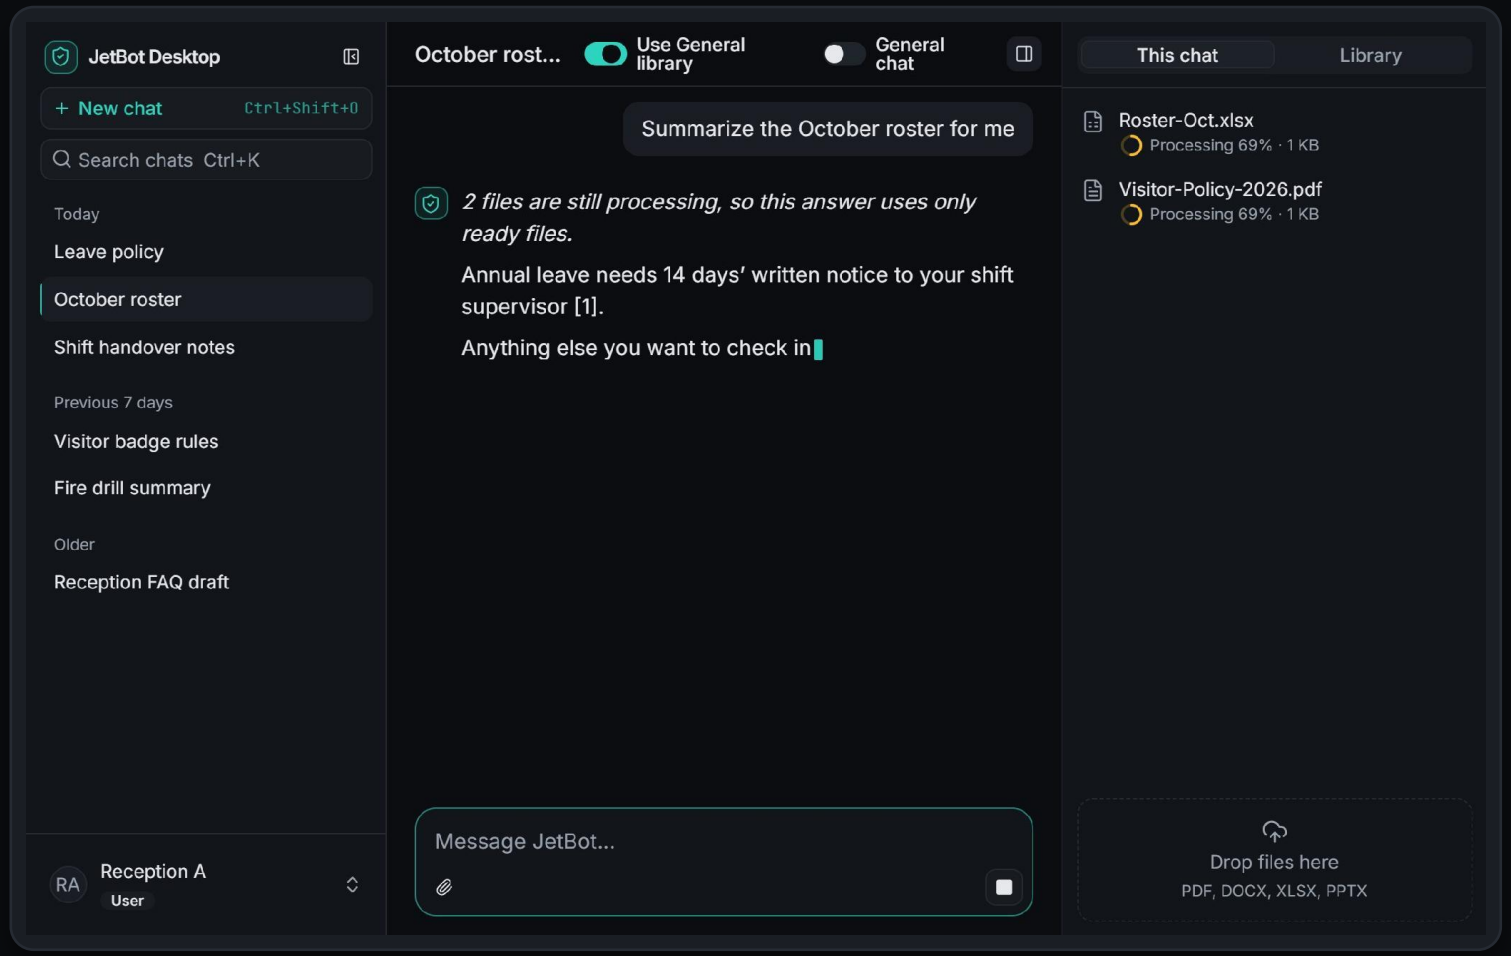
Accounts, shared library, models and a full audit log.



One installer. No server, no subscription, no internet after setup.

FEATURE · CHAT

Ask questions of your own files



- Attach PDF, Word, PowerPoint, Excel, CSV, HTML or text by drag-and-drop
- Files index in the background with live status: queued, processing, ready
- Each chat's files are private to that chat
- Follow-up questions keep context; long chats are summarised automatically
- Plain "General chat" mode when no documents are needed

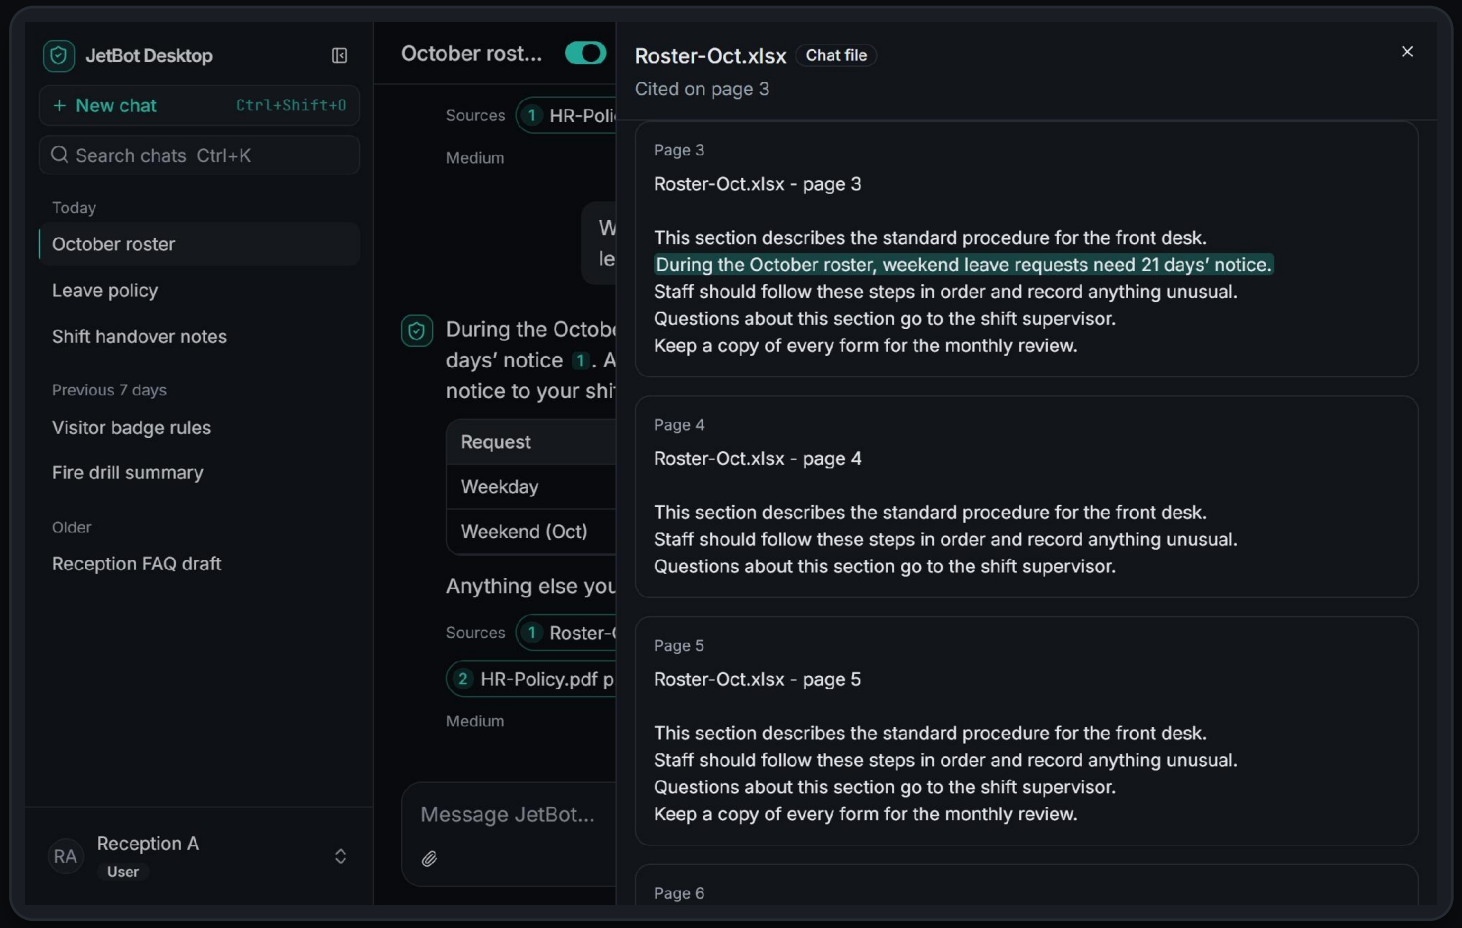
FEATURE · TRUST

Every answer shows its sources

Staff do not have to take the AI's word for it. Numbered citation chips show whether a fact came from the person's own file or the shared library, and one click opens the document at the cited page.

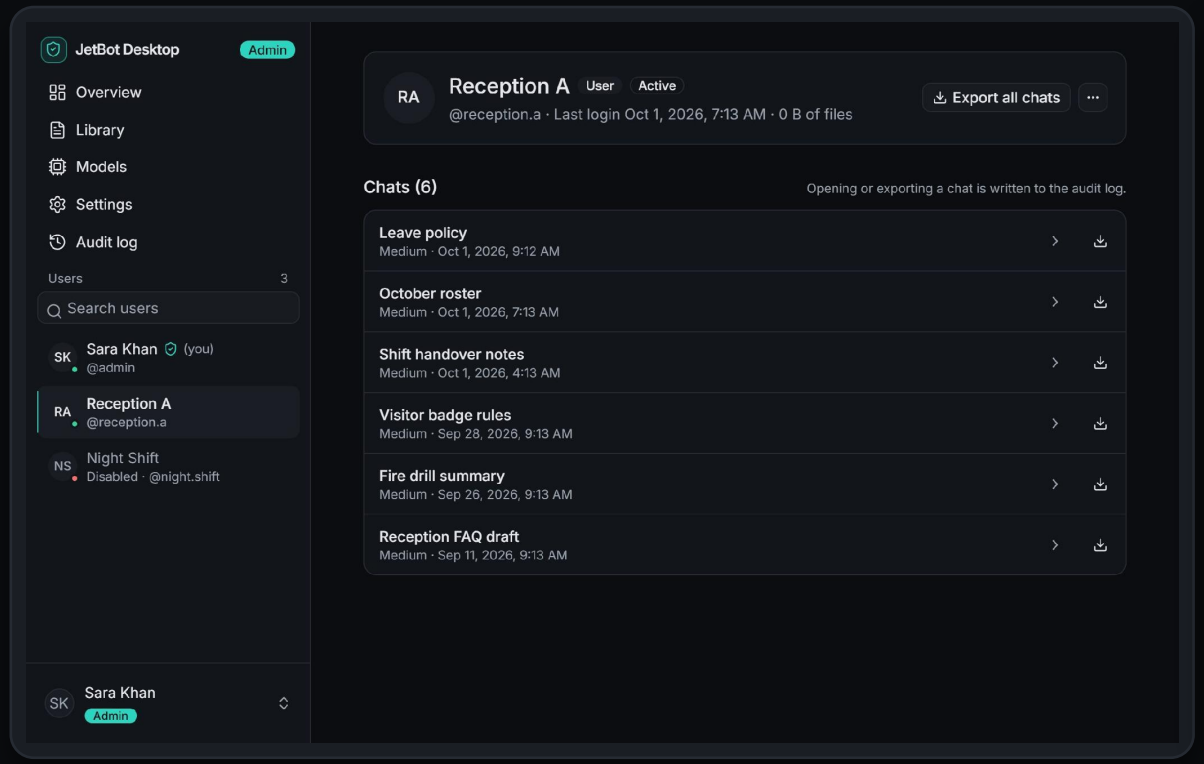
chat — the user's own file

General — the Admin's shared library

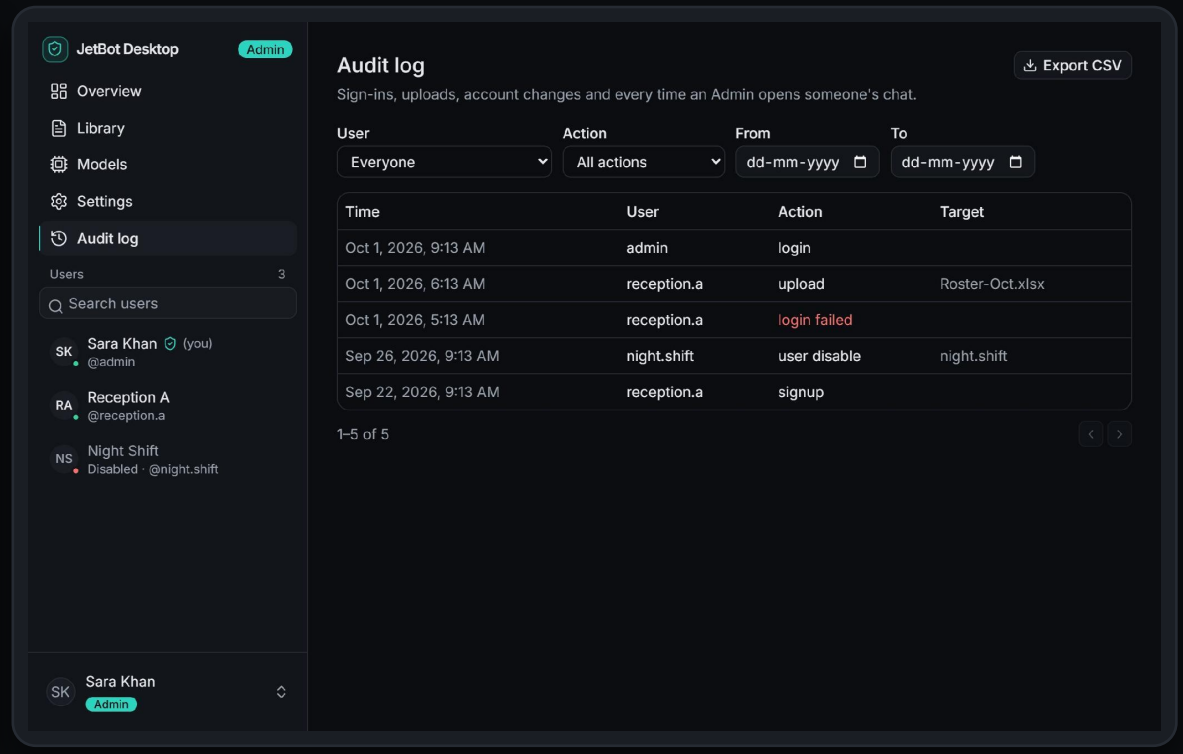


FEATURE · ADMINISTRATION

One Admin runs the device



Manage users, review and export chats as Markdown



Every sign-in, upload and chat review is audit-logged

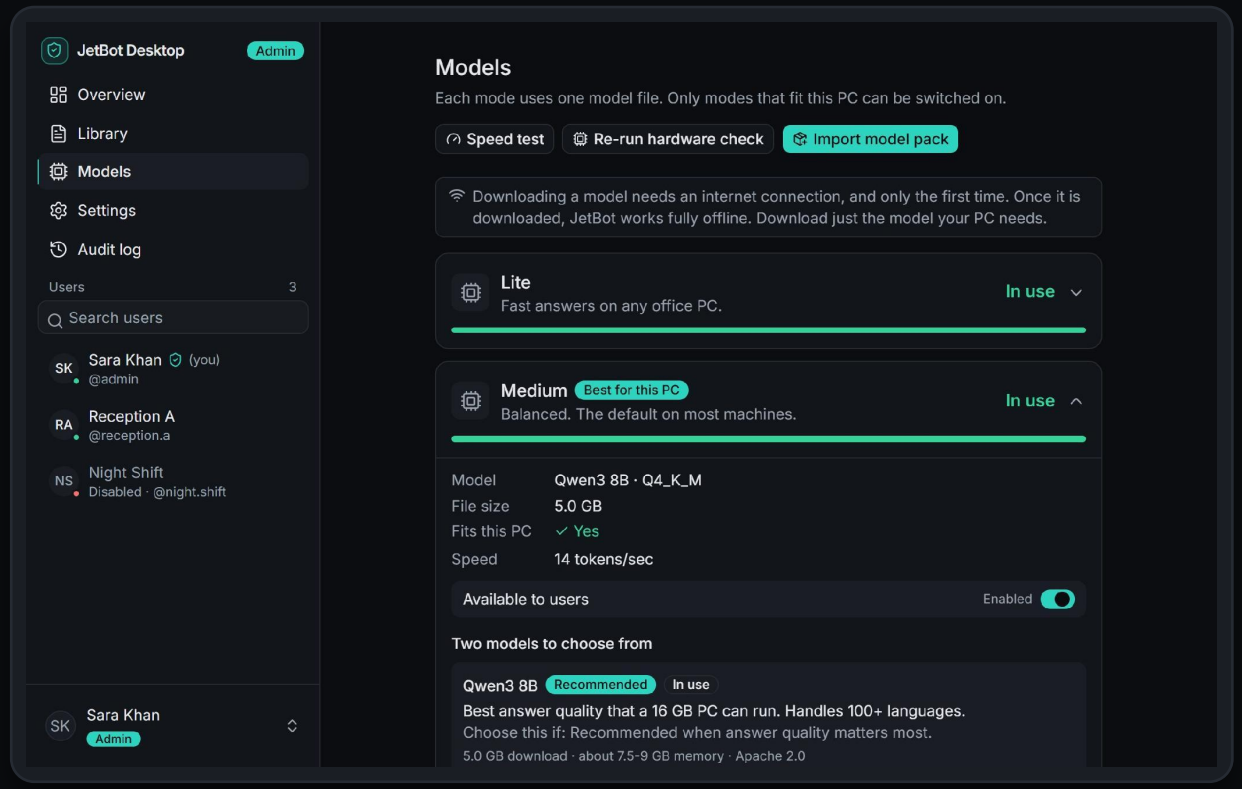
Plus: a shared General library, sign-up on/off, auto-lock, chat retention, storage quotas and offline updates.

FEATURE · MODELS AND GPU

The right model for the PC you have

Tier	Typical PC	Default model
Lite	8 GB RAM	Llama 3.2 3B
Medium	16 GB RAM or 6 GB GPU	Qwen3 8B
Advanced	32 GB RAM or 10–12 GB GPU	Qwen3 14B

- Automatic hardware check picks the tier on first run
- GPU acceleration today via Vulkan (NVIDIA, AMD, Intel Arc), with automatic CPU fallback
- Models arrive by in-app download or signed offline packs on a USB stick



SECURITY AND COMPLIANCE

Security designed in, not bolted on

Encrypted at rest

SQLCipher database, AES-256 files; key bound to the machine by Windows DPAPI.

No outbound traffic

The app can only talk to itself on 127.0.0.1, enforced by the content policy.

Strong sign-in

Argon2id passwords, lockout after 5 failures, one-time Admin recovery code.

Strict separation

Users never see each other's chats or files; Admins cannot chat.

Full audit trail

Sign-ins, uploads, account changes and every Admin chat view.

Signed and verified

Installer, models, packs and updates are signed and hash-checked at launch.

DEPLOYMENT

From USB stick to first answer in minutes

01

Install

Run the signed installer from a USB stick. No internet needed.

02

Set up

Create the Admin and save the one-time recovery code.

03

Check hardware

The app measures the PC and recommends a model tier.

04

Add knowledge

Admin uploads policies and guides to the General library.

05

Ask

Staff sign up and start chatting with their documents.

Installer 351 MB with offline WebView2. Data stays in C:\ProgramData and is kept on uninstall unless the Admin chooses to remove it. Updates are applied offline from a signed package, with an automatic backup first.

PERFORMANCE

Fast enough on an ordinary laptop

Measured on a low-end dev laptop: Intel i5-11400H, 7.6 GB RAM, CPU only, Qwen3 1.7B

~80 ms

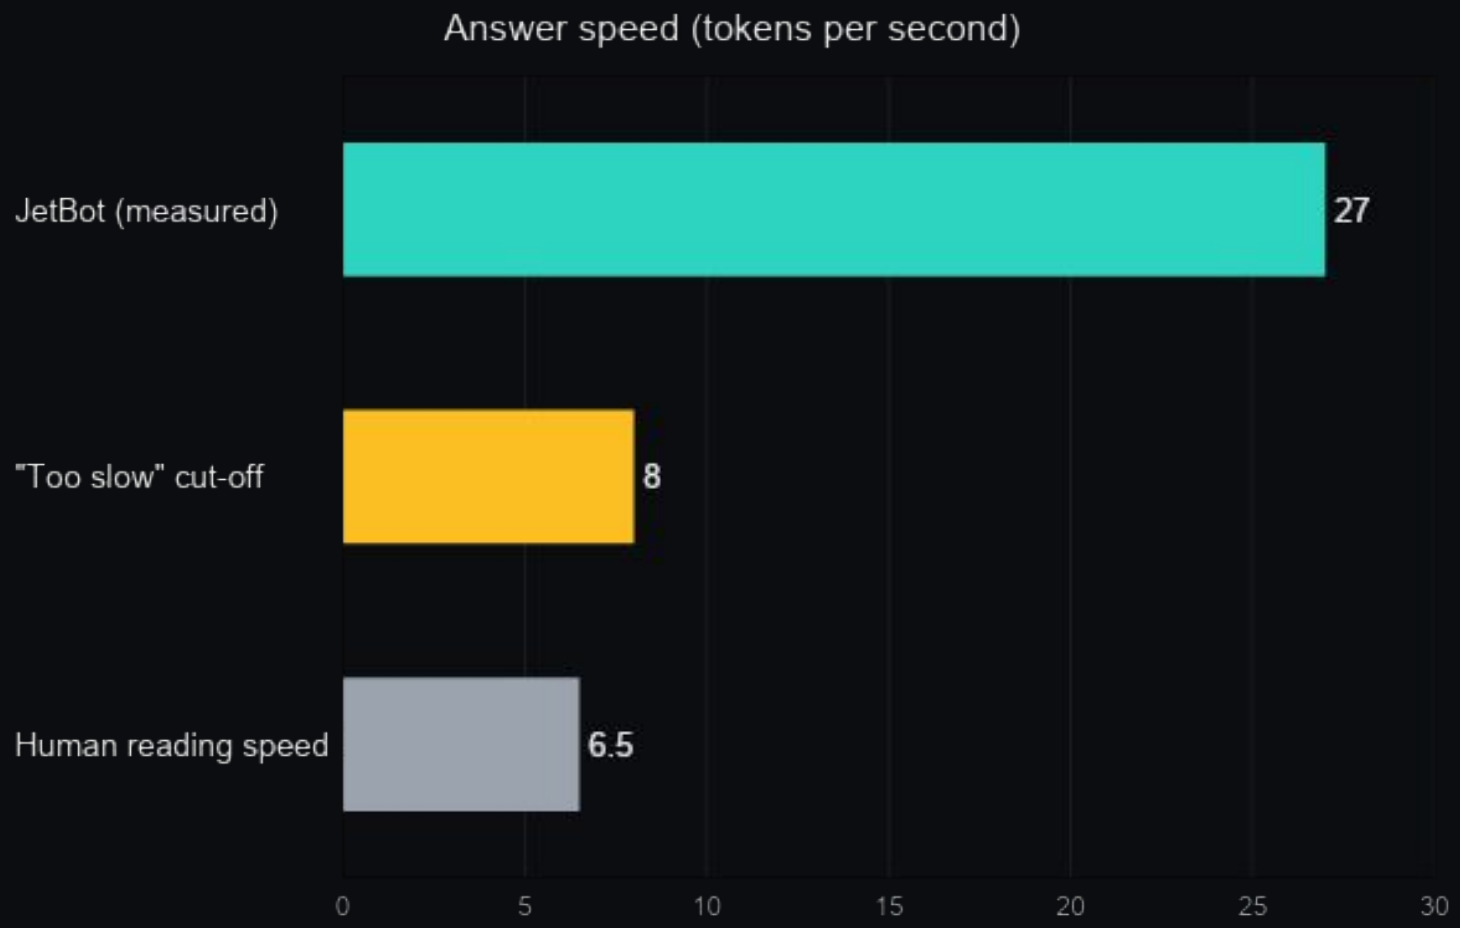
first word, warm

25–29

tokens per second, sustained

~10 s

app ready from launch



BUSINESS VALUE

Why customers buy it

0

documents sent to any cloud

Passes data-residency and confidentiality reviews that block cloud AI.

1

PC, many staff

Shared-shift devices serve every person on the rota with private accounts.

$0

per-user monthly fees

A one-off licence on hardware the customer already owns.

Target markets

- Defence and government
- Hospitals and clinics
- Law firms
- Banks and insurers
- Factories and field sites
- Front desks and call centres

ROADMAP

Where the product goes next

Today · v0.1

- Offline chat with files and citations
- General library and Admin console
- Lite, Medium, Advanced model tiers
- GPU acceleration (Vulkan) with CPU fallback
- Encryption, audit log, signed updates

Next · v1.x

- OCR for scanned PDFs and photos
- Multilingual search (Urdu, Arabic, Hindi and 100+ more)
- NVIDIA CUDA build for faster GPUs
- Smarter ranking of passages
- PDF export of chats, storage quota bar

Later · v2.x

- Vision: ask about charts, photos and screenshots
- Voice input and read-aloud, offline
- Active Directory and smart-card sign-in
- LAN mode: one GPU server, many desktops
- macOS and Linux builds

SUGGESTIONS

More ideas worth building

Features that raise the value per seat and open new markets

OCR and handwriting

Read scanned forms, faxes and photographed pages offline.

Compare documents

Show what changed between two versions of a policy or contract.

Watched folders

Keep the library in sync with a network share automatically.

Knowledge-gap analytics

Show the Admin which questions go unanswered, so docs can be fixed.

Templates and drafting

Fill letters, reports and forms from the library with one prompt.

Translate on device

Ask in Urdu or Arabic about English documents, and back.

Voice at the desk

Hands-free questions for reception and shop-floor staff.

Multi-site sync

Push signed library and model updates to many sites by USB or LAN.

Your knowledge stays here.

JetBot Desktop brings modern AI to the places the cloud cannot reach.

Book a pilot on your own hardware